

Glossary

p95 latency

p95 (95th percentile) latency is the response time that 95% of requests are faster than. It is commonly used to measure tail latency in systems. If p95 is 300ms, then 95% of requests complete within 300ms and 5% take longer.

p50 latency

Median response time (50th percentile).

RAG

Retrieval-Augmented Generation: retrieve relevant chunks from a knowledge base, then optionally use an LLM to generate an answer grounded in those chunks.

Embeddings

Vector representations of text used to measure semantic similarity.

FAISS

A library for fast similarity search over vectors (indexes).